

Universal simulation of Turing machine in $\mathcal{O}(T \log T)$ -time

Kenenbek Arzymatov

April 19, 2018

Introduction

Turing Machine.

Definition. Formally, a TM M is described by a tuple (Γ, Q, δ) containing:

- A finite set of the symbols that M 's tapes can contain. We assume that contains a designated “blank” symbol, denoted \square ; a designated “start” symbol, denoted \triangleright ; and the numbers 0 and 1. We call the alphabet of M .
- A finite set Q of possible states M 's register can be in. We assume that Q contains a designated start state, denoted q_{start} , and a designated halting state, denoted q_{halt} .
- A function $\delta : Q \times \Gamma^k \rightarrow Q \times \Gamma^{k-1} \times \{L, S, R\}^k$, where $k \geq 2$, describing the rules M use in performing each step. This function is called the transition function of M .

Statements:

1. Behavior of a Turing machine is determined by its transition function.
2. Every string in $\{0, 1\}^*$ represents some Turing machine.
3. Every TM is represented by infinitely many strings.

Definition. We denote by M_α the machine whose representation as a bit string is α . An algorithm (i.e., a machine) can be represented as a bit string once we decide on some canonical encoding.

Definition. The running time is the number of these basic operations performed. We measure it in asymptotic terms, so we say a machine runs in time $T(n)$ if it performs at most $T(n)$ basic operations time on inputs of length n .

Definition. Let $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$ and let $T : N \rightarrow N$ be some functions, and let M be a Turing machine. We say that M computes f if for every $x \in \{0, 1\}^*$, whenever M is initialized to the start configuration on input x , then it halts with $f(x)$ written on its output tape. We say M computes f in $T(n)$ -time if its computation on every input x requires at most $T(|x|)$ steps.

EFFICIENCY AND RUNNING TIME

Definition Computing a function and running time

Let $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$ and let $T : \mathbb{N} \rightarrow \mathbb{N}$ be some functions, and let M be a Turing machine. We say that M computes f if for every $x \in \{0, 1\}^*$, whenever M is initialized to the start configuration on input x , then it halts with $f(x)$ written on its output tape. We say M computes f in $T(n)$ -time if its computation on every input x requires at most $T(|x|)$ steps.

Time-constructible functions

A function $T : \mathbb{N} \rightarrow \mathbb{N}$ is time constructible if $T(n) \geq n$ and there is a TM M that computes the function $x \rightarrow T(|x|)$ in time $T(n)$. (As usual, $T(|x|)$ denotes the binary representation of the number $T(|x|)$.) Examples for time-constructible functions are $n, n \log n, n^2, 2n$. Almost all functions encountered in this book will be time constructible, and we will restrict our attention to time bounds of this form. (Allowing time bounds that are not time constructible can lead to anomalous results.) The restriction $T(n) \geq n$ is to allow the algorithm time to read its input.

Most of the specific details of our definition of Turing machines are quite arbitrary. It is a simple exercise to see that most changes to the definition do not yield a substantially different model, since our model can simulate any of these new models. In context of computational complexity, however, we have to verify not only that one model can simulate another, but also that it can do so efficiently. Now we state a few results of this type, which ultimately lead to the conclusion that the exact model is unimportant if we are willing to ignore polynomial factors in the running time. Variations on the model studied include restricting the alphabet Γ to be $\{0, 1, \square, \triangleright\}$, restricting the machine to have a single work tape, or allowing the tapes to be infinite in both directions. All results in this section are proved sketchily—completing these sketches into full proofs.

Claim

For every $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$ and time-constructible $T : \mathbb{N} \rightarrow \mathbb{N}$ if f is computable in time $T(n)$ by a TM M using alphabet Γ , then it is computable in time $4 \log |\Gamma| T(n)$ by a TM \tilde{M} using the alphabet $\{0, 1, \square, \triangleright\}$.

Short proof

Let M be a TM with alphabet Γ , k tapes, and state set Q that computes the function f in $T(n)$ time. We describe an equivalent TM \tilde{M} computing f with alphabet $\{0, 1, \square, \triangleright\}$, k tapes and a set Q' of states. The idea behind the transformation is simple: One can encode any member of Γ using $\log |\Gamma|$ bits. Thus, each of \tilde{M} 's work tapes will simply encode one of M 's tapes. For every cell in M 's tape we will have $\log |\Gamma|$ cells in the corresponding tape of \tilde{M} (see 1).

To simulate one step of M , the machine \tilde{M} will (1) use $\log |\Gamma|$ steps to read from each tape the $\log |\Gamma|$ bits encoding a symbol of Γ , (2) use its state register to store the symbols read, (3) use M 's transition function to compute the symbols M writes and M 's new state given this information, (4) store this information in its state register, and (5) use $\log |\Gamma|$ steps to write the encodings of these symbols on its tapes.




Figure 1: We can simulate a machine M using the alphabet $\{\square, \triangleright, a, b, \dots, z\}$ by a machine M' using $\{0, 1, \square, \triangleright\}$ via encoding every tape cell of M using five cells of M' .

One can verify that this can be carried out if \tilde{M} has access to registers that can store M 's state, k symbols in Γ , and a counter from 1 to $\log |\Gamma|$. Thus, there is such a machine \tilde{M} utilizing no more than $c|Q||\Gamma|^{k+1}$ states for some absolute constant c . (In general, we can always simulate several registers using one register with a larger state space. For example, we can simulate three registers taking values in the sets A, B and C , respectively, with one register taking a value in the set $A \times B \times C$, which is of size $|A||B||C|$.)

Oblivious Turing machines

Claim

Define a single-tape Turing machine to be a TM that has only one read-write tape, that is used as input, work, and output tape. For every $f : \{0, 1\}^* \rightarrow \{0, 1\}$ and time-constructible $T : \mathcal{N} \rightarrow \mathcal{N}$, if f is computable in time $T(n)$ by a TM M using k tapes, then it is computable in time $5kT(n)^2$ by a single-tape TM \tilde{M} .

Proof

Again the idea is simple: The TM \tilde{M} encodes k tapes of M on a single tape by using locations $1, k+1, 2k+1, \dots$ to encode the first tape, locations $2, k+2, 2k+2, \dots$ to encode the second tape etc. (see Figure 2). For every symbol a in M 's alphabet, \tilde{M} will contain both the symbol a and the symbol a^\wedge . In the encoding of each tape, exactly one symbol will be of the " \wedge type", indicating that the corresponding head of M is positioned in that location (see Figure 2). \tilde{M} will not touch the first $n+1$ locations of its tape (where the input is located) but rather start by taking $\mathcal{O}(n^2)$ steps to copy the input bit by bit into the rest of the tape, while encoding it in the above way.

To simulate one step of M , the machine \tilde{M} makes two sweeps of its work tape: First it sweeps the tape in the left-to-right direction and records to its register the k symbols that are marked by " \wedge ". Then \tilde{M} uses M 's transition function to determine the new state, symbols, and head movements and sweeps the tape back in the right-to-left direction to update the encoding accordingly. Clearly, \tilde{M} will have the same output as M . Also, since on n -length inputs M never reaches more than location $T(n)$ of any of its tapes, \tilde{M} will never need to reach more than location $2n + kT(n) \leq (k+2)T(n)$ of its work tape, meaning that for each of the at most $T(n)$ steps of M , \tilde{M} performs at most $5 \cdot k \cdot T(n)$ work (sweeping back and forth requires about $4 \cdot k \cdot T(n)$ steps, and some additional steps may be needed for updating head movement and book keeping). ■

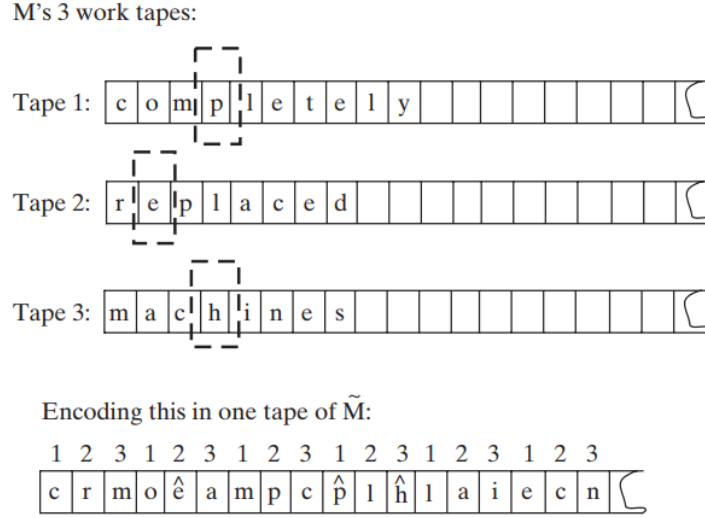


Figure 2: Simulating a machine M with three tapes using a machine \tilde{M} with a single tape.

With a bit of care, one can ensure that the proof of Claim 1.6 yields a TM \tilde{M} with the following property: Its head movements do not depend on the input but only depend on the input length. That is, every input $x \in \{0,1\}^*$ and $i \in \mathcal{N}$, the location of each of M 's heads at the i th step of execution on input x is only a function of $|x|$ and i . A machine with this property is called oblivious, i.e it has following

The universal Turing machine

Definition. There is a universal Turing machine U that can simulate any other Turing machine given its bit representation. Given a pair of bit strings (x, α) as input, this machine simulates the behavior of M_α on input x . This simulation is very efficient: If the running time of M_α was $T(|x|)$, then the running time of U is $O(T(|x|) \log T(|x|))$. Turing was the first to observe that general-purpose computers are possible, by showing a universal Turing machine that can simulate the execution of every other TM M given M 's description as input.

But it is good to remember why it was once counterintuitive. The parameters of the universal TM are fixed—alphabet size, number of states, and number of tapes. The corresponding parameters for the machine being simulated could be much larger. The reason this is the ability to use encodings. Even if the universal TM has a very simple alphabet, this suffices to represent the other machine's state and transition table on its tapes and then follow along in the computation step by step.

Efficient universal Turing machine

Claim

Theorem. There exists a TM U such that for every $x, \alpha \in \{0,1\}^*$, $U(x, \alpha) = M_\alpha(x)$, where M_α denotes the TM represented by α . Moreover, if M_α halts on input x within T steps then $U(x, \alpha)$ halts within $CT \log T$ steps, where C is a number independent

of $|x|$ and depending only on M_α 's alphabet size, number of tapes, and number of states.

Proof

We show a universal *TM* U such that given an input x and a description of a *TM* M that halts on x within T steps, U outputs $M(x)$ within $\mathcal{O}(T \log T)$ time (where the constants hidden in the \mathcal{O} notation may depend on the parameters of the *TM* M being simulated).

U will use its input and output tape in the same way M does and will also have extra work tapes to store M 's transition table and current state and to encode the contents of M 's work tapes.

Now we show a way to encode all of M 's work tapes in a single tape of U , which we call the main work tape of U . Let k be the number of tapes that M uses (apart from its input and output tapes) and its alphabet. We may assume that U uses the alphabet Γ^k (as this can be simulated with a overhead depending only on $k, |\Gamma|$). Thus we can encode in each cell of U 's main work tape k symbols of Γ , each corresponding to a symbol from one of M 's tapes. This means that we can think of U 's main work tape not as a single tape but rather as k parallel tapes.

Encoding M 's tapes on U 's tape

We encode the information (the content of M 's work tapes) using “buffer zones”: Rather than having each of U 's parallel tapes correspond exactly to a tape of M , we add a special kind of blank symbol \diamond to the alphabet of U 's parallel tapes with the semantics that this symbol is ignored in the simulation.

For convenience, we think of U 's parallel tapes as infinite in both the left and right directions (this can be easily simulated with minimal overhead: see Claim 1.8). Thus, we index their locations by $0, \pm 1, \pm 2, \dots$. Normally we keep U 's head on location 0 of these parallel tapes. We will only move it temporarily to perform a shift when, following our general approach, we simulate a left head movement by shifting the tape to the right and vice versa. At the end of the shift, we return the head to location 0.

We split each of U 's parallel tapes into zones that we denote by $R_0, L_0, R_1, L_1, \dots$. The cell at location 0 is not at any zone. Zone R_0 contains the two cells immediately to the right of location 0 (i.e., locations $+1$ and $+2$), while Zone R_1 contains the four cells $+3, +4, +5, +6$. Generally, for every $i \geq 1$, Zone R_i contains the $2 \cdot 2^i$ cells that are to the right of Zone R_{i-1} (i.e., locations $[2i+1-1, \dots, 2i+2-2]$). Similarly, Zone L_0 contains the two cells indexed by -1 and -2 , and generally Zone L_i contains the cells $[-2^{i+2}+2, \dots, -2^{i+1}+1]$. We shall always maintain the following invariants:

- Each of the zones is either empty, full, or half-full with non- \diamond is either $0, 2^i$, or $2 \cdot 2^i$ and the same holds number of symbols in zone R_i that are not for L_i . (We treat the ordinary \square symbol the same as any other symbol in Γ , and in particular a zone full of \square 's is considered full.) We assume that initially all the zones are half-full. We can ensure this by filling half of each zone with \diamond symbols in the first time we encounter it.
- The total number of non- \diamond symbols in $R_i \cup L_i$ is $2 \cdot 2^i$. That is, either R_i is empty and L_i is full, or R_i is full and L_i is empty, or they are both half-full.

- Location 0 always contains a non- \diamond symbol.

Performing a shift

The advantage in setting up these zones is that now when performing the shifts, we do not always have to move the entire tape, but we can restrict ourselves to only using some of the zones. We illustrate this by showing how U performs a left shift on the first of its parallel tapes (see also Figure 1.9):

1. U finds the smallest i_0 such that R_{i_0} is not empty. Note that this is also the smallest i_0 such that L_{i_0} is not full. We call this number i_0 the *index* of this particular shift.
2. U puts the leftmost non- \diamond symbol of R_{i_0} in position 0 and shifts the remaining leftmost $2^{i_0} - 1$ non- \diamond symbols from R_{i_0} into the zones R_0, \dots, R_{i_0-1} filling up exactly half the symbols of each zone. Note that there is exactly room to perform this since all the zones R_0, \dots, R_{i_0-1} were empty and indeed $2^{i_0} - 1 = \sum_{j=0}^{i_0-1} 2^j$.
3. U performs the symmetric operation to the left of position 0. That is, for j starting from $i_0 - 1$ down to 0, U iteratively moves the $2 \cdot 2^j$ symbols from L_{j+1} to fill half the cells of L_{j+1} . Finally, U moves the symbol originally in position 0 (modified appropriately according to M 's transition function) to L_0 .
4. At the end of the shift, all of the zones $R_0, L_0, \dots, R_{i_0-1}, L_{i_0-1}$ are half-full, R_{i_0} has 2^{i_0} fewer non- \diamond symbols, and L_i has 2^i additional non- \diamond symbols. Thus, our invariants are maintained.
5. The total cost of performing the shift is proportional to the total size of all the zones $R_0, L_0, R_{i_0}, L_{i_0}$. That is, $\mathcal{O}(\sum_{j=0}^{i_0} 2 \cdot 2^j) = \mathcal{O}(2^{i_0})$ operations.

After performing a shift with index i the zones $L_0, R_0, L_{i-1}, R_{i-1}$ are half-full, which means that it will take at least 2^{i-1} left shifts before the zones L_0, \dots, L_{i-1} become empty or at least 2^{i-1} right shifts before the zones R_0, \dots, R_{i-1} become empty. In any case, once we perform a shift with index i , the next 2^{i-1} shifts of that particular parallel tape will all have index less than i . This means that for every one of the parallel tapes, at most a $1/2^i$ fraction of the total number of shifts have index i . Since we perform at most T shifts, and the highest possible index is $\log T$, the total work spent in shifting U 's k parallel tapes in the course of simulating T steps of M is

$$\mathcal{O}(k \cdot \sum_{i=1}^{\log T} \frac{T}{2^{i-1}} 2^i) = \mathcal{O}(T \log T). \blacksquare$$